Machine Learning for Textual and Unstructured Data

Lecture 5:
Convolutional Neural Networks

Stephen Hansen
University College London

Image Classification

Image classification was one of the core problems that drove the growth of deep learning after 2010.

Problem is to correctly identify features of images that have been tagged by humans (e.g. ImageNet).

Image classification is now essentially a solved problem.

The main architecture for the problem is a *convolutional neural network*.

These have 'convolutional' and 'pooling' layers that reduce the dimensionality of images by associating them with learned local features.

Overall Idea [Chollet et al., 2022]

Figure 8.2 The visual world forms a spatial hierarchy of visual modules: elementary lines or textures combine into simple objects such as eyes or ears, which combine into high-level concepts such as "cat."

Example of Convolutional Filter [Chollet et al., 2022]

Original input

Single filter

Response map,
quantifying the presence
of the filter's pattern at
different locations

Figure 8.3 The concept of
a response map: A 2D map of
the presence of a pattern at
different locations in an input

Example of (Max) Pooling [James et al., 2021]

$$\text{Max pool} \begin{bmatrix} 1 & 2 & 5 & 3 \\ 3 & 0 & 1 & 2 \\ 2 & 1 & 3 & 4 \\ 1 & 1 & 2 & 0 \end{bmatrix} \rightarrow \begin{bmatrix} 3 & 5 \\ 2 & 4 \end{bmatrix}$$

Overall Architecture [James et al., 2021]

FIGURE 10.8. *Architecture of a deep CNN for the* **CIFAR100** *classification task. Convolution layers are interspersed with* 2×2 *max-pool layers, which reduce the size by a factor of* 2 *in both dimensions.*

Applications of CNNs

Economic Statistics

Measuring economic performance is a precursor for conducting effective economic policy.

Yet the state of measurement is imperfect, particularly in developing countries.

[Silungwe et al., 2022] report that

1. Half of countries do not publish quarterly GDP statistics from expenditure.
2. 25% of countries (and 50% of African countries) do not conduct a household budget survey.

The cost of increasing survey coverage is enormous, but unstructured data is plentiful.

Survey Coverage in Africa

A Consumption/income surveys   **B** Asset surveys

# of nationally-representative surveys, 2000–2010

4
3
2
1
0

Measuring Poverty from Satellite Data

[Jean et al., 2016] uses publicly available satellite data to measure poverty at local level in Africa.

Paper adopts a three-step process:

1. Begin with CNN trained on ImageNet data.
2. Update CNN to predict nightlight intensity, which is available at high resolution throughout the world.
3. Use learned features as input into ridge regression for predicting available local spending and asset surveys.

Nightlight intensity does a poor job at capturing variation in expenditure at low levels.

Local features of geography and buildings contain more information.

Filters for Predicting Nightlight Intensity

Fig. 2. Visualization of features. By column: Four different convolutional filters (which identify, from left to right, features corresponding to urban areas, nonurban areas, water, and roads) in the convolutional neural network model used for extracting features. Each filter "highlights" the parts of the image that activate it, shown in pink. By row: Original daytime satellite images from Google Static Maps, filter activation maps, and overlay of activation maps onto original images

Generalizability of Ridge Regression



A **Consumption expenditures**

Images in Media Data

[Ash et al., 2021] digitize both text and images from articles in the NYT and Fox News from 2000 through 2020.
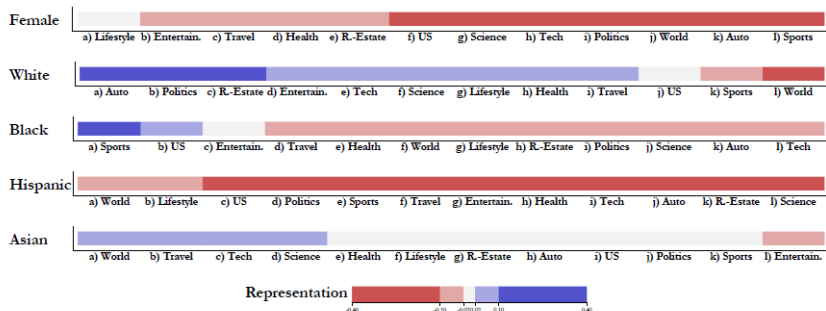
Use publicly available dataset of tagged images to build classifier for gender and ethnicity.

Convolutional neural network for image classification.

Representation in Media vs Population

Representation in Types of Articles

Association Between Images and Text

Other Applications

[Borgschulte et al., 2021] use CNN to predict age from CEO photos, correlate with economic distress.

[Mueller et al., 2021] use CNN to monitor local destruction in Syrian war.

References I

Ash, E., Durante, R., Grebenshchikova, M., and Schwarz, C. (2021).
Visual Representation and Stereotypes in News Media.

Borgschulte, M., Guenzel, M., Liu, C., and Malmendier, U. (2021).
CEO Stress, Aging, and Death.

Chollet, F., Kalinowski, T., and Allaire, J. J. (2022).
Deep Learning with R.
Manning, Shelter Island, NY, second edition.

James, G., Witten, D., Hastie, T., and Tibshirani, R. (2021).
An Introduction to Statistical Learning: With Applications in R.
Springer, New York NY, 2nd ed. 2021 edition edition.

Jean, N., Burke, M., Xie, M., Davis, W. M., Lobell, D. B., and Ermon, S. (2016).
Combining satellite imagery and machine learning to predict poverty.
Science, 353(6301):790–794.

Mueller, H., Groeger, A., Hersh, J., Matranga, A., and Serrat, J. (2021).
Monitoring war destruction from space using machine learning.
Proceedings of the National Academy of Sciences, 118(23):e2025400118.

References II

Silungwe, A., Baer, A., and Guerreiro, V. (2022).

2020 Global Stocktaking of National Accounts Statistics: Availability for Policy and Surveillance.

IMF Working Papers.